

Reading BabyLM scores: which 2025 numbers are meaningful, what the field achieved, and which evaluation bugs the 2026 pipeline fixed

Working note

July 19, 2026

Abstract

This is a self-contained guide to interpreting BabyLM challenge scores, aimed at someone preparing a 2026 entry who wants to know three things: which of the many evaluation numbers actually carry signal, how a given score compares to the field, and which of the concrete evaluation bugs reported in the 2025 papers have been fixed in the 2026 pipeline. The findings are drawn from a full reading of all 41 papers of the 2025 proceedings (`2025.babylm-main`) and from reading the actual source of the 2026 evaluation pipeline (`github.com/babylm-org/babylm-eval`). **Headline conclusions.** (i) Only two scores are trustworthy for ranking: **BLiMP** and **(Super)GLUE**; the rest are at chance, at the noise floor, or were outright buggy. (ii) BLiMP spreads widely across systems (≈ 49 – 80) and is the axis that separates models; (Super)GLUE is compressed (≈ 53 – 72) and barely reflects a $10\times$ difference in training data. (iii) The 2026 pipeline fixes the two most damaging scoring bugs (BLiMP tie/NaN inflation; age-of-acquisition mis-fit) and deletes the unreliable WUG tasks, but the EWoK “concepts-absent-from-training” problem remains.

Contents

1	Purpose, method, and how to read this	2
2	The BabyLM setting (2025 and 2026)	2
3	How a submission is scored	3
4	Which scores are meaningful?	3
4.1	The reliable core: BLiMP and (Super)GLUE	3
4.2	Real but too noisy to rank: Entity Tracking	4
4.3	Broken at this scale: EWoK (and COMPS)	4
4.4	Noise floor or buggy: the human-likeness half	5
4.5	Why the aggregate is not a safe ranking	6
5	The 2026 pipeline: which evaluation bugs were fixed	6
6	What the BLiMP and (Super)GLUE spread looks like	7
6.1	Reference bands for placing your own score	8

7	Recommendations for a 2026 entry	9
8	Bottom line	9
A	Paper number $\rightarrow$ identity	9
B	Selected evidence (load-bearing quotes and numbers)	10

1 Purpose, method, and how to read this

The BabyLM Challenge asks entrants to pretrain a language model on a human-scale amount of text — 100M words or fewer (Strict), or 10M or fewer (Strict-Small) — and evaluates it on a fixed suite of linguistic and cognitive benchmarks. Because that suite reports on the order of a dozen numbers per model, and because those numbers are then blended into a single leaderboard rank, it is easy to over-read a score that is really noise. This report exists to prevent that.

It answers three questions. **(1) Which scores are meaningful?** (Section 4) — i.e. reliable across seeds, discriminative between models, construct-valid, and not stuck at chance. **(2) What does the field look like?** (Section 6) — the actual spread of the two trustworthy metrics, with reference bands so a new score can be placed without the leaderboard. **(3) Which 2025 bugs did 2026 fix?** (Section 5) — audited against the real pipeline code, not a changelog (there is none).

Method and sourcing. The reliability verdicts come from a full reading of the 41 papers in *Proceedings of the First BabyLM Workshop* (2025.babylm-main.1 through .41); the organizers’ summary is .28. Individual papers are cited by anthology number, e.g. (m.16) for 2025.babylm-main.16; a full number $\rightarrow$ identity map is in Appendix A, and a condensed set of the load-bearing quotes and numbers is in Appendix B. The bug-fix status was established by cloning and diffing the 2025 (babylm/evaluation-pipeline-2025) and 2026 (babylm-org/babylm-eval) pipelines. This report is intended to be read on its own.

2 The BabyLM setting (2025 and 2026)

Tracks and budgets. Table 1 summarizes the 2026 structure. The core Strict / Strict-Small text tracks are unchanged from 2025; the standalone Interaction and Multimodal tracks of 2025 were folded back into Strict (teacher feedback and multimodal data are now allowed within the word budget rather than being their own competitions); and a new Multilingual track was added.

Track	Budget	Notes
Strict	$\leq 100\text{M}$ words	or up to 1B with the ≤ 10 -epoch exposure cap
Strict-Small	$\leq 10\text{M}$ words	or up to 100M with the exposure cap
Multilingual (new 2026)	100M tokens	English/Dutch/Chinese, byte-premium adjusted

Table 1: 2026 competition tracks. All tracks cap training exposure at ≤ 10 epochs and require intermediate checkpoints (every 1M words to 10M, every 10M to 100M). A detoxified training corpus is provided.

Evaluation philosophy. Most of the suite is *zero-shot*: the model assigns (pseudo-)log-probabilities to candidate strings, and accuracy is whether it prefers the correct one, exactly as in BLiMP min-

imal pairs. The organizers investigated replacing the finetuned (Super)GLUE tasks with zero-shot prompting but concluded that robust in-context learning does not emerge at this data scale (induction heads require billions of tokens), so finetuning was kept. This philosophy carries into 2026, which is why the 2025 reliability picture below is directly relevant to a 2026 entry.

3 How a submission is scored

Every text-track model is reduced to two headline numbers and one aggregate.

- **NLP score** — an *accuracy* average over the forced-choice / labelled tasks: BLiMP, BLiMP-Supplement, EWok, Entity Tracking, COMPS, and the finetuned (Super)GLUE suite (BoolQ, MNLI, MRPC, MultiRC, QQP, RTE, WSC; MNLI/QQP subsampled to 10k; CoLA, SST2, MNLI-mm, QNLI dropped as redundant).
- **Human-likeness score** — a *correlation-to-humans* average over reading-time prediction (surprisal vs human reading times), the two WUG morphology tasks (model inflection preference vs the human distribution), and age-of-acquisition (model word-learning curves vs WordBank).
- **Macro Average** — the leaderboard sort key.

The aggregation is a two-level unweighted mean: each task score is the mean of its subtasks, and the macro is the mean across tasks. Checking against the organizers’ results table pins the exact structure down: **Macro = mean(NLP score, Human-likeness score)** — a flat 50/50 split (e.g. CLASS-IT 20.4/52.9→36.6; MoEP 31.5/53.2→42.3). *This 50/50 split is the crux of Section 4*: half the flagship number comes from the least reliable half of the suite. (Note: some online summaries quote a 50%/30%/20% BLiMP/GLUE/MSGs weighting — that is the 2023 scheme; MSGs was not used in 2025.)

What the field actually achieved. Table 2 gives the 2025 track winners against the strongest provided baseline. Two facts stand out and inform how to read a rank. First, human-likeness and NLP were awarded separately because a system could win either; and at the top of the largest track (Strict-Small) the two winners are near-opposites (MoEP 31.5/53.2 vs AMLM 8.4/58.3), i.e. the two halves are decoupled. Second, in the Strict track *no submission beat the strongest GPT-BERT baseline on the macro average* — the awards go to the best non-baseline entries.

4 Which scores are meaningful?

Table 3 is the summary; the subsections give the evidence. “At chance” means ≈ 0.50 for a two-alternative task. The verdicts are unusually well-supported because they are *convergent across many independent submissions* and, in the strongest cases, stated by the winning teams and diagnosed mechanistically.

4.1 The reliable core: BLiMP and (Super)GLUE

BLiMP is the anchor. It rises smoothly with training — “BLiMP scores demonstrate a gradual and consistent increase during training” (m.34) — and the Strict-Small NLP winner’s own audit, an appendix literally titled “*Should we trust BabyLM Metrics?*”, concludes that for BLiMP and (Super)GLUE the answer “appears to be yes... Failed and untrained models perform expectedly poorly on these” (m.31). Three caveats keep it from being read at fine granularity: a scoring artifact

Track	System (award)	Human-lik.	NLP	Macro
Strict	Simple-Diffusion (NLP winner)	12.6	58.4	35.5
	CLASS-IT (human-lik. winner)	20.4	52.9	36.6
	<i>GPT-BERT baseline (best)</i>	<i>24.2</i>	<i>62.7</i>	<i>43.5</i>
Strict-Small	AMLM-Hard-Decay (NLP winner)	8.4	58.3	33.3
	MoEP (human-lik. winner)	31.5	53.2	42.3
	<i>GPT-BERT baseline (best)</i>	<i>18.7</i>	<i>56.1</i>	<i>37.4</i>
Interaction	BLM / Once Upon a Time (both)	20.8	54.4	37.6

Table 2: 2025 track winners vs the best provided baseline, from the organizers’ results table. Note the Strict baseline is not beaten on Macro, and the Strict-Small winners are anti-aligned across the two headline metrics.

could inflate degenerate models in 2025 (Section 5); it collapses toward chance under weak compute (“BLiMP... hovers 44–52 across configurations,” m.12; 48.7 for a heavily-quantized model, m.11); and it saturates at the top, where the best models approach the human/large-LM reference.

The **BLiMP-Supplement** is weaker: several models sit near 0.50 on it (m.5), and one paper traced a supplement “advantage” entirely to three subcategories containing speaker labels that only one condition’s preprocessing preserved — a headline difference that was a data-cleaning artifact, not model quality (m.22).

(**Super**)**GLUE** is the trustworthy secondary: accuracy-based, moves with model quality (m.31), with MNLI/MRPC/RTE carrying the real signal (m.1). Its caveats are cost and patchy sensitivity: it requires finetuning (the expensive, higher-variance stage); it barely reflects the data budget (in Figure 1 the best Strict-Small 10M-word systems sit at or above most Strict 100M-word systems on GLUE); it is sometimes wholly inert (frozen at 57.7 across every configuration in m.36; ≈ 63.5 flat in m.41); and WSC/RTE carry 4–12-point seed variance, enough that m.19 dropped WSC from its analysis. Read the aggregate, not the individual small tasks.

4.2 Real but too noisy to rank: Entity Tracking

Entity Tracking is the most *method-sensitive* accuracy task — multi-token prediction (+4–5 pts, m.41), contrastive data (+7%, m.2), and recombination (12.95→19.65, m.40) all move it — but it is also the highest-variance metric on the board: seed SD of 6–9 points (20.70 ± 6.09 ; 31.68 ± 8.75 , m.22), mid-run collapse (“rises to 41.8 at 20M tokens before collapsing to 13.4,” m.34; “drops sharply after the first epoch,” m.41), and a barely-trained model beating a strong GPT-BERT on it (m.36). Use it only as a seeded diagnostic for a specific intervention, never as a ranking component.

4.3 Broken at this scale: EWoK (and COMPS)

EWoK is at chance in essentially every text-track paper that reports it: 52.7–56.7 across all sizes (m.1); 53.0–53.8 across all methods (m.2); 0.494–0.519 across 14 curricula (m.26); “none of the models perform better than a random guess” (m.41). The outstanding paper m.15 supplies the mechanism, and it is a *data* problem: EWoK “demonstrates no sensitivity to changes in architecture or training strategy, with performance remaining around 50%...”, because in 37.7% of EWoK items at least one tested concept appears fewer than 100 times in the training corpus and in 13% both

Task	Type	Verdict	Core evidence
BLiMP	accuracy	Trust	Monotone with training (m.34); separates models; but tie artifact (m.16), chance under weak compute (m.11,m.12), saturating at top
(Super)GLUE	accuracy (FT)	Trust	Trusted by the winners’ audit (m.31); but sometimes inert (m.36,m.41), high seed variance on RTE/WSC (m.19); costly
BLiMP-Supplement	accuracy	Probe only	Often near chance (m.5); a supplement “gap” was a speaker-label preprocessing artifact (m.22)
Entity Tracking	accuracy	Probe only	Seed SD 6–9 pts (m.22); mid-run collapse 41.8→13.4 (m.34); U-shaped; tokenizer-sensitive
COMPS	accuracy	Noise	At chance for weak models (m.14,m.22); floor-limited
EWoK	accuracy	Broken	≈0.50 everywhere; 13% of items’ concepts never appear in training (m.15)
Reading-time	correlation	Noise	≈0, model-insensitive (m.31,m.32); tokenizer-fragile (m.41)
WUG past-tense	correlation	Broken	Eval bug (m.14); sign-flipping negatives (m.22,m.38)
WUG adjective	correlation	Broken	Untrained model scores ≈78 (m.31); non-discriminative (m.32)
Age-of-Acquisition	correlation	Broken	Sigmoid fit on 1–5 points → “misleading” (m.29); decided an award (m.39)

Table 3: Per-task reliability verdict for the 2025 suite. **Trust** = reliable enough to rank on; **Probe only** = real signal but too noisy/format-sensitive to rank on; **Noise/Broken** = do not read small differences.

concepts appear zero times — “the training dataset does not properly support EWoK evaluation.” Where EWoK *does* move it is suspect: a morpheme tokenizer pushes it to 67–71, which the authors themselves call “surprising, as EWoK measures basic world knowledge rather than a linguistic task” (m.21), i.e. a segmentation/log-probability artifact. **COMPS** shares the floor problem (at chance for weak models, m.14,m.22) though strong baselines reach 56–60. Do not optimize either, and do not read small differences.

4.4 Noise floor or buggy: the human-likeness half

Reading-time prediction has near-zero magnitude and is insensitive to model quality: the winners’ audit found the numbers “stay relatively similar, regardless of the model” (m.31); “the correlations of all models are below 0.1” (m.32). Values sit at 0.0–0.1 partial- R^2 , flip sign when the tokenizer changes (m.41), and can perversely favour *smaller* models (m.30,m.17).

WUG (past-tense and adjective-nomination) is unreliable in several independent ways. There was an outright evaluation bug — “a bug was discovered in the evaluation for the WuG task... we therefore exclude this task” (m.14). Under a morpheme tokenizer every configuration scores *exactly* 100.00, a normalization artifact (m.21). The past-tense correlation routinely goes strongly

negative and sign-flips across near-identical models (-20.7 vs $+37.5$, m.14; -24 to $+2$ with SD up to ± 10 , m.22). And the adjective task is gameable: “one could obtain quite high scores on adjective nominalization, around 78, just from initialization, no training required” (m.31).

Age-of-Acquisition was, in 2025, the most dangerous number on the board — it was *miscomputed*. m.29 inspected the pipeline: “only very few data points (1–5 words) are considered... due to an unpassed condition on the parameters of the fitted sigmoid function ... Limited data points lead to either a score of zero or a strong negative correlation; hence, these results can be misleading.” This explains the wild spread (mostly -0.07 to 0 with lone outliers, m.34; literal 0.00 at $p=1.0$, m.40; -79.6 , m.29). It was also inflated by broken/untrained runs (m.31). *This specific bug is fixed in 2026* (Section 5).

4.5 Why the aggregate is not a safe ranking

The individual-task problems would matter less if they averaged out. They do not, because the macro weights the broken half at 50% and the noisiest tasks have the widest ranges. The decisive demonstration (m.31): a team accidentally trained a model with too-low a batch size, spiking the loss; it “performs expectedly poorly on... BLiMP and (Super)GLUE” but “remarkably well on... EWoK, adjective nominalization, COMPS, and AoA,” and its aggregate *beat every properly-trained model*. Their verdict: “the final scoring metric seems to be unfairly skewed.” Relatedly, one real award (Strict-Small human-likeness, MoEP, m.39) turned almost entirely on the miscomputed AoA task: its AoA (53.70) was $4\text{--}14\times$ any baseline’s, while on BLiMP it was *below* every baseline. Run-to-run macro SD is $\approx 1\text{--}2$ points (m.14) — larger than many ranking gaps — and cross-metric non-transfer is pervasive (grammar and knowledge/reasoning move independently or oppositely under the same intervention: m.1,m.2,m.5,m.11,m.12), so no single scalar is honest. Collapsing the metrics “would obscure these distinctions and reduce interpretability” (m.13).

5 The 2026 pipeline: which evaluation bugs were fixed

Table 4 is the summary; the notes give the mechanism and code evidence, established by reading the 2026 source and its git history.

BLiMP tie/NaN (fixed). In 2025, `rank_and_evaluate` chose the argmax candidate with `torch.max(...)[1]`, which returns the *first* index on ties and lets NaNs create spurious maxima; an order-invariant or degenerate model could therefore be scored as correct, and a dummy model could report 100.0. The 2026 code first maps NaN to $-\infty$ and then breaks ties *randomly*, so a model that cannot distinguish the minimal pair now scores at chance. It does *not* remove the 22 paradigms whose two sentences are permutations of one another — they are still in the data snapshot; they simply no longer inflate a degenerate model.

Age-of-Acquisition (fixed). The 2025 model-side sigmoid fit was unbounded and capped at `maxfev=1000`; on many clean learning curves the shape parameter diverged and the word was discarded, leaving the correlation computed over a handful of survivors — exactly the “unpassed condition on the parameters of the fitted sigmoid” that m.29 diagnosed. The 2026 fit is *bounded*, uses $20\times$ more evaluations, averages surprisals within each checkpoint, corrects the chance-surprisal ceiling, and gates on a p -value. One caveat: the 2026 strict README reports AoA as an MSE-like quantity computed server-side, so the exact headline number may differ from the repository’s Pearson-correlation code; only the repo-side fit could be verified here.

WUG (removed). Rather than repair the morphological WUG tasks, 2026 deletes them — the scoring routine, the task dispatch, and the collation entries are all gone. The “WUG=100” artifact

Issue (reported by)	2025 behaviour	2026 state	Verdict
BLiMP tie / NaN inflation (m.16)	<code>torch.max(...)[1]</code> picks first index on ties; NaNs propagate; a degenerate model can report 100.0	NaN $\rightarrow -\infty$ and <i>random</i> tie-breaking; degenerate models now score at chance (9bcf1f2, Apr 2026)	Fixed
— underlying permutation items	22 of 67 BLiMP paradigms are word-order permutations	Same 67 paradigms still shipped; only the <i>scoring</i> changed	Partial
Age-of-Acquisition mis-fit (m.29)	Unbounded sigmoid fit (<code>maxfev=1000</code>) diverges and drops most words, so the correlation uses only 1–5 points	Bounded fit, <code>maxfev=20000</code> , per-checkpoint averaging, corrected chance ceiling, <i>p</i> -value gate (7d0a6c8, Jul 2026)	Fixed
WUG bug (m.14); WUG=100 under morpheme tokenizers (m.21)	Buggy morphological WUG (past/adj) correlation tasks	Both WUG tasks <i>deleted</i> from the 2026 suite (9bcf1f2)	Removed
EWoK at chance: ~13% of items use concepts absent from training (m.15)	Filter only drops out-of-vocab items; no training-coverage check	<code>dl_and_filter.py</code> and <code>vocab.txt</code> byte-identical; pair construction/collation fixed but coverage untouched	Not fixed
Reading-time fragility (m.17,m.31)	Surprisal-RT regression	Unchanged except a tokenizer-loading fallback	—

Table 4: Status in the 2026 pipeline (`babylm-org/babylm-eval`) of evaluation bugs reported in the 2025 papers.

and the sign-flipping correlations are therefore moot, but so is any WUG signal.

EWoK (not fixed). The reason EWoK sits at chance is that $\approx 13\%$ of its items use concepts absent from the training corpus (m.15) — a property of the *data*, and the 2026 filter and vocabulary files are byte-for-byte identical to 2025. Recent 2026 commits fixed how EWoK pairs are *constructed and scored* (two context-based minimal pairs; a prediction/target matching fix), which improves correctness, but the coverage problem is untouched. Treat EWoK the same in 2026 as in 2025: near-uninformative at this scale.

Other changes worth knowing. The 2026 strict zero-shot suite keeps BLiMP, BLiMP-Supplement, EWoK, Entity Tracking, COMPS, reading-time and AoA, *drops* both WUG tasks, and *adds* GlobalPIQA (length-normalized). Fine-tuning keeps the same seven (Super)GLUE tasks. A new Multilingual track (English/Dutch/Chinese) adds MultiBLiMP, Hanzi structure/pinyin minimal pairs, and MECO reading-time. The Macro Average is computed server-side in both years, so the exact 2026 weighting cannot be read from the pipeline code; but the *composition* of the two halves changed — the human-likeness side lost WUG, the accuracy side gained GlobalPIQA.

6 What the BLiMP and (Super)GLUE spread looks like

Three things stand out, and they are what matter for judging a new score.

BLiMP is the metric that separates systems. Its range here is about 31 points (49 to 80),




Figure 1: BLiMP versus (Super)GLUE for a cross-section of 2025 systems. **(A)** Each point is one system; colour is the track; the star is the official GPT-BERT baseline. **(B)** The same values as a one-dimensional spread per metric (black bar = mean, whisker = min-max). BLiMP ranges ≈ 49 –80; (Super)GLUE is compressed into ≈ 53 –72. Across systems the two are only weakly related (Pearson $r \approx 0.25$). Data compiled from the primary submission papers; each point’s BLiMP and GLUE come from the *same* source, but values are as-reported (including author re-implementations of the baselines) and GLUE-vs-SuperGLUE averaging conventions vary slightly. This is not the official leaderboard table, which is gated.

standard deviation near 9.5. A one-point BLiMP difference low down is meaningful; near the top it is not, because the strongest models approach the reference and the metric saturates.

(Super)GLUE is compressed and scale-insensitive. Its range is about 19 points (53 to 72), standard deviation near 6. In Figure 1 the best Strict-Small (10M-word) systems match or exceed most Strict (100M-word) systems on GLUE — a metric that barely reflects a $10\times$ data difference cannot finely rank models that differ by far less. Read GLUE as a coarse band, not a ranking.

The two are weakly correlated, so “good” is metric-specific. Systems with strong BLiMP and mediocre GLUE (Once Upon a Time: 77.5/53) and the reverse (Batchwise-Convergent: 49/58) both exist. A single system can look strong on one axis and ordinary on the other; the baseline (GPT-BERT) is the rare point at the top of both.

6.1 Reference bands for placing your own score

Table 5 gives approximate, track-dependent anchors that let a new score be judged without the leaderboard. To decide whether a submission is good, compare its **BLiMP** against the GPT-BERT baseline for its track (≈ 80 for Strict, high-60s to low-70s for Strict-Small) — that is the number that most cleanly says “ahead of” or “behind” the field — and treat **GLUE** as a wide competitive band. Do *not* use EWoK or the human-likeness tasks to decide whether a model is good; small movements there are noise.

Metric	Chance / floor	Weak	Competitive	Strong / reference
BLiMP	50 (minimal pairs)	50–62	67–75	76–80 (GPT-BERT); ≈ 88 human
(Super)GLUE	≈ 50 (majority)	53–57	58–66	67–72 (GPT-BERT ≈ 71)
EWoK	50	≈ 50 –54 for <i>everyone</i> — do not read differences		
Entity Track.	≈ 20	high-variance (seed SD 6–9 pts); use only as a seeded probe		

Table 5: Rough reference bands for the metrics that carry signal, from the 2025 cross-section. The BLiMP “human” anchor is the original benchmark’s human agreement ($\approx 88.6\%$); the organizers additionally plot a Llama-70B skyline that the best Strict/Interaction systems approach on BLiMP.

7 Recommendations for a 2026 entry

1. **Optimize BLiMP + BLiMP-Supplement first**, via the levers that actually moved them in 2025: training objective and architecture (diffusion MLM m.38, adaptive masking m.31, GPT-BERT-style dual objective, multi-token prediction m.41). It is the most reliable, most monotone, cheapest (zero-shot) signal.
2. **Treat GLUE as the secondary target**; budget for finetuning variance; report the aggregate, not WSC/RTE/MultiRC individually.
3. **Do not optimize EWoK, COMPS, or the human-likeness tasks for their own sake.** If you compete on human-likeness, note the award can be captured by a targeted push on one idiosyncratic task, but any single such score is fragile.
4. **Use Entity Tracking as a diagnostic, not a target**, and only with multiple seeds and confidence intervals.
5. **Report per-task, with seeds and confidence intervals**, following the best-practice examples in the field: m.2 (10 seeds, bootstrap CIs) and m.22 (4 seeds, per-task SDs). A single-seed macro number is not evidence that one model is better than another.
6. **Verify the pipeline before trusting a number** — the BLiMP tie-handling and AoA fit were both buggy in 2025 and both are worth re-checking in the 2026 code you actually run.

8 Bottom line

The trustworthy scores are BLiMP (read coarsely, compared to the track baseline) and (Super)GLUE (a coarse, compressed band). EWoK is at chance, the human-likeness tasks are at the noise floor or were buggy, and the 2025 Macro Average was unsafe enough that a broken training run topped it. The 2026 pipeline is genuinely cleaner: the two most damaging scoring bugs (BLiMP tie/NaN, AoA mis-fit) are fixed and the unreliable WUG tasks are gone; the one substantive survivor is EWoK’s data-coverage problem. Judge a submission by its BLiMP relative to the baseline, corroborate with GLUE, and read everything else with the scepticism the 2025 submissions themselves recommended.

A Paper number $\rightarrow$ identity

Numbers are ACL Anthology 2025.babylm-main.N. Only papers cited above are listed.

m.1 Velasco & Roque, *Rethinking the Role of Text Complexity* • **m.2** Ulm et al., *Contrastive Decoding for Synthetic Data* (10-seed CIs) • **m.5** Takmaz et al., *Model Merging* (multimodal) • **m.11** Aman et al.,

BitMar (low-bit multimodal) • **m.12** Loáiciga et al., *Smaller batch sizes for ELC-BERT* • **m.13** Gao et al., *BLiSS* (L2 selective tolerance) • **m.14** Haller et al., *BLaLM* (linear attention; reports the WUG bug) • **m.15** Ganescu et al., *Looking to Learn* (outstanding; EWoK data-coverage) • **m.16** Păpușoi & Nisioi, *Elementary Baselines* (BLiMP tie artifact) • **m.17** McCurdy et al., *Hall of Mirrors* (surprisal instability) • **m.19** Roque & Velasco, *Text Simplification + Curriculum* • **m.21** Bölücü & Can, *Morpheme-Aware* (tokenizer; WUG=100 artifact) • **m.22** *Do Syntactic Categories Help...* (4-seed variance; supplement artifact) • **m.26** Schoenegger et al., *Influence-driven Curriculum* • **m.28** *Findings of the Third BabyLM Challenge* (organizers) • **m.29** Padovani et al., *Dialogue Is Not Enough* (AoA miscomputation) • **m.31** Edman & Fraser, *Mask and You Shall Receive* / AMLM (SS NLP winner; “Should we trust BabyLM Metrics?”) • **m.32** Martins et al., *Once Upon a Time* / BLM (Interaction winner) • **m.33** Kamzela et al., *LLM-designed study plans* • **m.34** Fysikoudi et al., *Active Curriculum LM* • **m.36** *Batch-wise Convergent Pre-training* (GLUE frozen at 57.7) • **m.38** Kosmopoulou et al., *Masked Diffusion* (Strict NLP winner) • **m.39** Tapaninaho, *MoEP* (SS human-likeness winner; AoA-driven) • **m.40** Tampier et al., *RecombiText* • **m.41** Aynedinov & Akbik, *Multi-Token Prediction*.

B Selected evidence (load-bearing quotes and numbers)

BLiMP tie artifact (m.16). “In 22 of the 67 BLiMP subtasks, the two sentences... are permutations of the same multiset of words... ties are counted as correct, which inflates accuracy.” Removing the 22 subtasks drops a word-frequency baseline from 0.663 to 0.498 (chance).

Winners’ audit (m.31, Appendix “Should we trust BabyLM Metrics?”). An untrained model scores ≈ 78 on adjective-nomination; a broken run (spiked loss) beat every properly-trained model on the aggregate via EWoK/adj-nom/COMPS/AoA. Verdict: trust BLiMP and (Super)GLUE; reading/eye-tracking “stay relatively similar, regardless of the model”; “the rest... a baby fist of salt.”

AoA miscomputation (m.29). “only very few data points (1–5 words) are considered... due to an unpassed condition on the parameters of the fitted sigmoid function... these results can be misleading.”

EWoK data coverage (m.15). “EWoK... remaining around 50% regardless of the experiment conditions”; 37.7% of items have a concept appearing <100 times in training, 13% have both at zero — “the training dataset does not properly support EWoK evaluation.”

Entity-tracking instability (m.34, m.22). “rises to 41.8 at 20M tokens before collapsing to 13.4 by the end of training”; seed SD 6–9 points (20.70 ± 6.09).

GLUE insensitivity (m.36). (Super)GLUE frozen at exactly 57.7 across every proposed and controlled configuration. In the compiled cross-section (Fig. 1) the best 10M-word systems match or exceed most 100M-word systems on GLUE, i.e. it does not track the data budget.

Metric decoupling (m.32, m.13). RL “primarily enhances narrative coherence... while leaving surface-level fluency — measured by the BabyLM tasks — largely unchanged”; combining metrics into one score “would obscure these distinctions and reduce interpretability.”