

Logistic regression, step by step

Notes to go with the lecture slides

These notes go through the logistic regression lecture slowly. They follow the order of the slides and fill in the steps that the slides skip. Section 2, on the 0–1 loss, is the longest, because it is where the rest of the lecture starts. If a formula on the slides looked like it came out of nowhere, its derivation should be here.

You need only a little algebra: adding up a list of numbers, the exponential function e^z , and the natural logarithm $\log$. The rules for e and $\log$ that we use are collected in Appendix A.

Contents

1	What training means	2
2	Counting mistakes: the 0–1 loss	3
2.1	Two conditions	3
2.2	Choose the threshold 0	3
2.3	From two conditions to one: multiply by the label	3
2.4	A mistake is $y_i f(x_i) \leq 0$	4
2.5	The indicator function	4
2.6	Adding up: the 0–1 loss	4
2.7	The signed score $m = yf(x)$	5
2.8	Why not minimize the 0–1 loss directly?	5
3	The SVM as a loss: the hinge	6
4	Same score, another loss: the log loss	6
5	From score to probability: the sigmoid	7
6	The log loss is $-\log p(\text{true class})$	8
7	The score is the log-odds	9

8	Making predictions	9
9	Summary	10
A	Rules for e and $\log$ used in these notes	11

1 What training means

The data

We have n training examples. Example number i has an input x_i and a label y_i . In this lecture there are two classes, and we write the labels as numbers:

$$y_i = +1 \text{ (the positive class)} \quad \text{or} \quad y_i = -1 \text{ (the negative class)}.$$

The labels could have been called “pass” and “fail”, or 1 and 0. We use +1 and -1 because it makes the formulas in Section 2 short. We write the whole training set as $D = \{(x_1, y_1), \dots, (x_n, y_n)\}$, or $D = \{(x_i, y_i)\}$ for short.

The model

A model is a formula that turns an input into a prediction. It has some numbers in it, the parameters, that training has to choose. We collect the parameters in one symbol, θ (theta), and write the prediction for example i as

$$\hat{y}_i = f(x_i \mid \theta).$$

The hat on $\hat{y}_i$ means “predicted”, as opposed to the true label y_i . The bar in $f(x_i \mid \theta)$ is read “given”: the prediction depends on the input, once the parameters are fixed.

The model for the whole lecture is the linear score

$$f(x) = w \cdot x + b.$$

With one input feature, $w \cdot x$ is just w times x . With several features $x = (x_1, \dots, x_d)$ it is $w_1 x_1 + \dots + w_d x_d$. The parameters are $\theta = (w, b)$. This is the same score the linear SVM uses.

The loss

A loss is a number that says how badly a choice of parameters fits the training data. Training means finding the parameters with the smallest loss:

$$\min_{\theta} \mathcal{L}(\theta; D).$$

Read this as: “over all choices of θ , make $\mathcal{L}$ as small as possible.” Methods differ in their model and in their loss. Linear regression, for example, adds up squared errors,

$$\mathcal{L} = \sum_i (y_i - \hat{y}_i)^2.$$

The symbol $\sum_i$ means “add up over all examples $i = 1, \dots, n$ ”. With three examples, $\sum_i (y_i - \hat{y}_i)^2$ is $(y_1 - \hat{y}_1)^2 + (y_2 - \hat{y}_2)^2 + (y_3 - \hat{y}_3)^2$.

The question for this lecture is: what loss should a classifier use?

2 Counting mistakes: the 0–1 loss

The most natural loss for a classifier is the number of training examples it gets wrong. This section builds a formula for that number, one step at a time.

2.1 Two conditions

A score $f(x)$ is turned into a class by comparing it with a threshold t : predict +1 when the score is above t , and -1 when it is below. So we want

$$\begin{aligned} f(x_i) &> t && \text{for every example with } y_i = +1, \\ f(x_i) &< t && \text{for every example with } y_i = -1. \end{aligned}$$

Positives should score above the threshold, negatives below it. This is the same picture as in the ROC lecture. Each example either meets its condition (a correct prediction) or breaks it (a mistake). The four possibilities are the four cells of the table on the slide:

	predicted +1	predicted -1
true label +1	correct	mistake
true label -1	mistake	correct

The training error is the number of examples that break their condition. To write it as a formula we first turn the two conditions into one.

2.2 Choose the threshold 0

For a score like $w \cdot x + b$ the usual threshold is $t = 0$. This loses nothing: a threshold $t \neq 0$ can be moved into the score by changing b to $b - t$, because $w \cdot x + b > t$ says the same as $w \cdot x + (b - t) > 0$. (If f is a probability, the usual threshold is 0.5. We will see in Section 8 that for logistic regression “probability above 0.5” and “score above 0” are the same thing.)

With $t = 0$ the two conditions read

$$\begin{aligned} f(x_i) &> 0 && \text{if } y_i = +1, \\ f(x_i) &< 0 && \text{if } y_i = -1. \end{aligned}$$

2.3 From two conditions to one: multiply by the label

Here is the step that the slide does in one line. Look at the product $y_i f(x_i)$ of the label and the score.

- If $y_i = +1$, then $y_i f(x_i) = f(x_i)$. The condition $f(x_i) > 0$ is the same as $y_i f(x_i) > 0$.
- If $y_i = -1$, then $y_i f(x_i) = -f(x_i)$. The condition $f(x_i) < 0$ is the same as $-f(x_i) > 0$, because multiplying both sides of an inequality by -1 flips it. So it too is the same as $y_i f(x_i) > 0$.

Both conditions have become one:

$$y_i f(x_i) > 0 \quad \text{means example } i \text{ is classified correctly.}$$

The table below checks all four cases. The sign of the score is the prediction; the sign of the product says whether the prediction was right.

label y_i	sign of the score $f(x_i)$	sign of $y_i f(x_i)$	result
+1	+	+	correct
+1	−	−	mistake
−1	−	+	correct
−1	+	−	mistake

This works only because we chose the labels $+1$ and -1 . Multiplying by $+1$ leaves the score alone; multiplying by -1 flips its sign, which turns “should be negative” into “should be positive”.

2.4 A mistake is $y_i f(x_i) \leq 0$

The opposite of $y_i f(x_i) > 0$ is $y_i f(x_i) \leq 0$. So example i is a mistake when $y_i f(x_i) \leq 0$. The case $y_i f(x_i) = 0$ means the score is exactly 0: the point sits on the boundary and the classifier cannot decide. Counting it as a mistake is a convention.

2.5 The indicator function

To count mistakes with a formula, we need something that is 1 for a mistake and 0 for a correct example. That is the indicator function. For any statement A ,

$$\mathbf{1}[A] = \begin{cases} 1 & \text{if } A \text{ is true,} \\ 0 & \text{if } A \text{ is false.} \end{cases}$$

Some examples: $\mathbf{1}[3 > 2] = 1$, $\mathbf{1}[2 > 3] = 0$, $\mathbf{1}[-0.5 \leq 0] = 1$, $\mathbf{1}[1.5 \leq 0] = 0$. Some books write it as $\mathbb{1}[A]$ or $I(A)$; it is the same thing.

2.6 Adding up: the 0–1 loss

The cost of example i is $\mathbf{1}[y_i f(x_i) \leq 0]$: 1 for a mistake, 0 for a correct prediction. Adding up over all examples counts the mistakes:

$$\mathcal{L}_{0-1} = \sum_i \mathbf{1}[y_i f(x_i) \leq 0] = \text{number of training mistakes.}$$

It is called the 0–1 loss because each example costs either 0 or 1. Dividing by n gives the fraction of mistakes, the training error rate.

Worked example. Five training examples with these labels and scores:

i	y_i	$f(x_i)$	$y_i f(x_i)$	correct?	$\mathbf{1}[y_i f(x_i) \leq 0]$
1	+1	+2.0	+2.0	yes	0
2	+1	-0.5	-0.5	no	1
3	-1	-1.5	+1.5	yes	0
4	-1	+0.8	-0.8	no	1
5	+1	+0.3	+0.3	yes	0
$\mathcal{L}_{0-1}$ (sum of the last column)					2

Example 3 shows the trick at work: its score is negative, but so is its label, so the product is positive and the prediction is correct. Two of the five examples are mistakes, an error rate of 40%. We will reuse these five examples for the other two losses.

2.7 The signed score $m = yf(x)$

The product $yf(x)$ will appear in every loss from now on, so it gets a name here: the signed score of an example, $m = yf(x)$. It is the score, with its sign flipped for negative examples. (Many books call $yf(x)$ the margin. We avoid that word here, because in the SVM lecture the margin is the width of the band around the line.) Its sign says whether the example is classified correctly. Its size says how far the score is from the boundary: a large positive m is a confident correct prediction, a large negative m is a confident mistake.

On the slides, the loss plots use $m = yf(x)$ as the horizontal axis. As a function of m , the 0–1 loss is a step: 1 for $m \leq 0$ and 0 for $m > 0$ (the orange line in Figure 1).

2.8 Why not minimize the 0–1 loss directly?

It is the loss we care about, so why not train with it? There are three reasons.

1. It gives no direction. A small change to w or b usually changes no prediction at all, so the count stays the same; then, when a point crosses the boundary, it jumps by 1. A method that improves the parameters by following the slope of the loss (next week: gradient descent) finds no slope to follow. For example, in the hours-of-study data on the slides (Section 8), the fitted boundary is at 5.40 hours and the nearest training points are at 5.3 and 5.6 hours. Any boundary between those two gives the same 6 mistakes.
2. It is hard to minimize. Finding the line with the fewest training mistakes is NP-hard in general, as the number of features grows (Ben-David, Eiron and Long, 2003): no efficient method is known.
3. It has no preference among lines with the same count. A line that barely separates the classes and one with a wide margin cost the same.

So we use a stand-in: a loss that is easy to minimize and still punishes mistakes. The SVM uses the hinge loss, logistic regression uses the log loss.

3 The SVM as a loss: the hinge

The hinge loss of one example with signed score $m = yf(x)$ is

$$\max(0, 1 - m) = \begin{cases} 0 & \text{if } m \geq 1, \\ 1 - m & \text{if } m < 1. \end{cases}$$

Here $\max(a, b)$ is the larger of the two numbers. An example with signed score at least 1 pays nothing. Below 1, the cost grows by 1 for each unit that m drops: $m = 0.5$ costs 0.5, $m = 0$ costs 1, $m = -1$ costs 2, $m = -2$ costs 3. So mistakes cost more the further they are on the wrong side, and even correct examples pay a little if they are too close to the boundary ($0 < m < 1$).

The soft-margin SVM minimizes

$$\mathcal{L} = \frac{1}{2}\|w\|^2 + C \sum_i \max(0, 1 - y_i f(x_i)).$$

The first term keeps w small, which makes the margin wide; $\|w\|^2 = w_1^2 + \dots + w_d^2$. The second term is the total hinge loss, and C sets how much it counts compared with the first. At the solution, the slack ξ_i of the SVM lecture equals the hinge loss of example i , so this is the same problem as $\frac{1}{2}\|w\|^2 + C \sum_i \xi_i$.

Worked example. For the five examples above, the hinge losses are 0, 1.5, 0, 1.8 and 0.7, a total of 4.0. Example 5 is classified correctly ($m = 0.3$) but still pays 0.7, because it is inside the margin.

4 Same score, another loss: the log loss

Logistic regression replaces the hinge with the log loss,

$$\log(1 + e^{-m}), \quad m = yf(x).$$

Here $\log$ is the natural logarithm. Some values:

signed score m	-2	-1	0	0.5	1	2	4
0-1 loss	1	1	1	0	0	0	0
hinge loss	3	2	1	0.5	0	0	0
log loss	2.127	1.313	0.693	0.474	0.313	0.127	0.018

Three things to notice, comparing with the hinge (Figure 1):

- The log loss is never exactly zero. Even a confidently correct example pays a little, so every example has some say in where the boundary goes, not only the ones near it.
- Far on the wrong side it grows linearly, like the hinge. For very negative m , e^{-m} is huge, the 1 hardly matters, and $\log(1 + e^{-m}) \approx \log(e^{-m}) = -m$. For example, at $m = -5$ the log loss is 5.007.

- It is smooth: no corner at $m = 1$.




Figure 1: The three losses as functions of the signed score $m = yf(x)$. Positive m means a correct prediction. Unlike the 0–1 loss, the hinge and the log loss have a slope that points towards fewer mistakes, and both are convex (bowl-shaped), which makes them easy to minimize.

Logistic regression minimizes the same kind of objective as the SVM, with the log loss in place of the hinge:

$$\mathcal{L} = \frac{1}{2}\|w\|^2 + C \sum_i \log(1 + e^{-y_i f(x_i)}).$$

This is what scikit-learn’s `LogisticRegression` minimizes (it does not penalize b).

Worked example. For the five examples, the log losses are 0.127, 0.974, 0.201, 1.171 and 0.554, a total of 3.028.

So far the log loss is just another curve. The next two sections show where it comes from.

5 From score to probability: the sigmoid

The sigmoid function squashes any number into the interval between 0 and 1:

$$\sigma(z) = \frac{1}{1 + e^{-z}}.$$

At $z = 0$, $e^0 = 1$ and $\sigma(0) = 1/2$. For large positive z , e^{-z} is close to 0 and $\sigma(z)$ is close to 1. For large negative z , e^{-z} is huge and $\sigma(z)$ is close to 0. In between it rises smoothly: $\sigma(-2) = 0.12$, $\sigma(0) = 0.5$, $\sigma(2) = 0.88$.

Logistic regression reads the score as a probability:

$$p(y = +1 \mid x) = \sigma(f(x)) = \frac{1}{1 + e^{-(w \cdot x + b)}}.$$

The score is the same as before. The sigmoid changes how we read it, not where the boundary is.

The probability of the true class is $\sigma(yf)$

Write f for the score $f(x)$. The probability of the negative class is what is left over:

$$\begin{aligned}
 p(y = -1 \mid x) &= 1 - \sigma(f) \\
 &= 1 - \frac{1}{1 + e^{-f}} = \frac{(1 + e^{-f}) - 1}{1 + e^{-f}} = \frac{e^{-f}}{1 + e^{-f}} && \text{(common denominator)} \\
 &= \frac{e^{-f} \cdot e^f}{(1 + e^{-f}) \cdot e^f} = \frac{1}{e^f + 1} && \text{(multiply top and bottom by } e^f) \\
 &= \sigma(-f) && \text{(the definition of } \sigma \text{ with } z = -f).
 \end{aligned}$$

So $p(y = +1 \mid x) = \sigma(f)$ and $p(y = -1 \mid x) = \sigma(-f)$. This is the same trick as in Section 2: multiplying by the label y covers both cases at once,

$$p(\text{true class of } x) = \sigma(yf(x)) = \sigma(m).$$

For the five examples, the probabilities the model gives to the true class are 0.881, 0.378, 0.818, 0.310 and 0.574. The two mistakes (examples 2 and 4) get less than 0.5.

6 The log loss is $-\log p(\text{true class})$

A good model should give high probability to what actually happened. So a natural cost for one example is $-\log p(\text{true class})$. It is 0 when the model gives the true class probability 1, and it grows without limit as that probability goes to 0:

$p(\text{true class})$	0.99	0.9	0.5	0.1	0.01
$-\log p(\text{true class})$	0.01	0.11	0.69	2.30	4.61

Confident and wrong is expensive. Now put in $p(\text{true class}) = \sigma(yf)$ from the last section:

$$\begin{aligned}
 -\log p(\text{true class}) &= -\log \sigma(yf) \\
 &= -\log \frac{1}{1 + e^{-yf}} && \text{(the definition of } \sigma) \\
 &= \log(1 + e^{-yf}) && \text{(because } -\log(1/a) = \log a).
 \end{aligned}$$

This is exactly the log loss of Section 4. So minimizing the log loss means: choose w and b so that the model gives high probability to the true class of every training example. For the five examples, $-\log 0.881 = 0.127$, $-\log 0.378 = 0.974$, and so on: the same numbers as the log losses we computed directly.

7 The score is the log-odds

Run the sigmoid backwards: if the model says probability p , which score z produced it? Solve $p = \sigma(z)$ for z :

$$\begin{aligned}
 p &= \frac{1}{1 + e^{-z}} \\
 1 + e^{-z} &= \frac{1}{p} && \text{(flip both sides)} \\
 e^{-z} &= \frac{1}{p} - 1 = \frac{1-p}{p} && \text{(subtract 1)} \\
 -z &= \log \frac{1-p}{p} && \text{(take log of both sides)} \\
 z &= \log \frac{p}{1-p} && \text{(because } -\log a = \log(1/a)\text{)}.
 \end{aligned}$$

The ratio $p/(1-p)$ is the odds: how many times more likely the positive class is than the negative one. A probability of 0.8 is odds of $0.8/0.2 = 4$, “4 to 1”. Its logarithm is the log-odds, also called the logit. So

$$w \cdot x + b = \log \frac{p}{1-p} :$$

logistic regression models the log-odds as a linear function of x . That is where “regression” in its name comes from, even though it is a classifier.

This gives the score a meaning. A score of 0 is odds of 1, even odds, $p = 0.5$. Each extra unit of score multiplies the odds by $e \approx 2.72$; a score of 1 is $p = 0.731$. And each extra unit of an input x_j adds w_j to the score, so it multiplies the odds by e^{w_j} .

8 Making predictions

To predict a class, threshold the probability. The default is 0.5:

$$p > 0.5 \iff \sigma(f(x)) > 0.5 \iff f(x) > 0.$$

The middle step holds because σ is increasing and $\sigma(0) = 0.5$. So the boundary is where the score is 0, the hyperplane $w \cdot x + b = 0$: a point with one feature, a line with two, a plane with three. This is the same kind of boundary as the linear SVM’s. The two methods usually draw different hyperplanes on the same data, because they minimize different losses: every example pulls on the logistic regression line, only the support vectors on the SVM’s.

Any other threshold t works the same way. By Section 7,

$$p > t \iff f(x) > \log \frac{t}{1-t}.$$

Each threshold gives one point on the ROC curve, as in the ROC lecture.

The hours-of-study example. On the slides, 20 students (simulated data), with hours of study as the only feature and pass as the positive class. Fitted with the log loss alone (no $\frac{1}{2}\|w\|^2$ term), the model is

$$f(x) = 0.668x - 3.611, \quad p(\text{pass} \mid x) = \sigma(0.668x - 3.611).$$

- A student who studied 3 hours has score -1.61 and $p(\text{pass}) = 0.17$. One who studied 7 hours has score $+1.07$ and $p(\text{pass}) = 0.74$, odds of 2.91.
- The boundary $p = 0.5$ is where $0.668x - 3.611 = 0$, at $x = 3.611/0.668 = 5.40$ hours. With this threshold the model makes 6 mistakes on the 20 students (an error rate of 30%): three who passed with fewer than 5.40 hours, and three who failed with more.
- Each extra hour multiplies the odds of passing by $e^{0.668} = 1.95$, so it roughly doubles them. For example, the odds go from 2.91 at 7 hours to 5.67 at 8 hours.
- With threshold $t = 0.3$ the score must exceed $\log(0.3/0.7) = -0.847$, which happens above 4.14 hours. With $t = 0.7$ it must exceed $+0.847$, above 6.67 hours.

9 Summary

loss	cost of one example	notes
0-1	$\mathbf{1}[yf(x) \leq 0]$	counts mistakes; what we care about, but flat, hard to minimize
hinge (SVM)	$\max(0, 1 - yf(x))$	zero beyond the margin; only support vectors matter
log (logistic regression)	$\log(1 + e^{-yf(x)})$	equals $-\log p(\text{true class})$; smooth; never zero

- With labels ± 1 and threshold 0, “correct” is the single condition $yf(x) > 0$.
- Logistic regression and the linear SVM use the same score $f(x) = w \cdot x + b$ and both draw a hyperplane. They differ in the loss.
- Logistic regression reads the score as a probability, $p(y = +1 \mid x) = \sigma(f(x))$; the score itself is the log-odds.
- A logistic regression is one artificial neuron: a weighted sum of the inputs plus b , through the sigmoid. Next week we connect many of them, and see how to find their weights.

A Rules for e and $\log$ used in these notes

$\log$ is the natural logarithm, the inverse of e^z : $\log(e^z) = z$ and $e^{\log a} = a$ for $a > 0$. Here $e \approx 2.718$.

rule	example of use
$e^0 = 1$	$\sigma(0) = 1/(1 + 1) = 1/2$
$e^a e^b = e^{a+b}$, so $e^{-f} e^f = 1$	$p(y = -1) = 1/(e^f + 1)$
$\log(1/a) = -\log a$	$-\log \sigma(yf) = \log(1 + e^{-yf})$
$e^{z+1} = e^z \cdot e$	each unit of score multiplies the odds e^z by e
$-\log a = \log(1/a)$	solving $p = \sigma(z)$ for z
$\log 1 = 0$	a probability of 1 costs $-\log 1 = 0$
$e^z > 0$ for every z	so $0 < \sigma(z) < 1$

Reference. S. Ben-David, N. Eiron and P. M. Long, “On the difficulty of approximately maximizing agreements”, *Journal of Computer and System Sciences* 66 (2003) 496–514.